

# 1 Overview


- **The puzzle:** How do recurrent neural networks (RNNs) use continuous vectors to represent discrete symbolic structures?
- **Findings:**
  - RNNs trained on structure-dependent tasks learn to implicitly implement *tensor product representations*.
  - Several popular tasks for training sentence encoders are not structure-sensitive enough to induce RNNs to capture sentence structure.

# 2 Tensor Product Representations

- A principled method for representing compositional symbolic structures in vector space (Smolensky 1990)
- Represent the input with pairs of **fillers** and **roles**:  
Cats chase dogs = **cats**:subject + **chase**:verb + **dogs**:object
- Each filler  $f_i$  and role  $r_i$  has a vector embedding.
- The representation of the input is the sum of the outer products of each  $f_i$  and  $r_i$ :  $\sum f_i \otimes r_i$

# 3 Tensor Product Decomposition

- **Goal:** Approximate an RNN's learned encodings (such as E above) with a tensor product representation.
- **Approach:** (right, top) Train a model to generate tensor product representations that are close to the RNN's encodings.
  - **Loss:** Mean squared error between  $E'$  and E
- **Evaluation:** (right, bottom) Pass this model's output,  $E'$ , to the RNN's decoder.



# 4 Role Schemes

	3	1	1	6
Left-to-right	0	1	2	3
Right-to-left	3	2	1	0
Bidirectional	(0,3)	(1,2)	(2,1)	(3,0)
Wickelroles	#_1	3_1	1_6	1_#
Tree	L	RLL	RLR	RR
Bag-of-words	$r_0$	$r_0$	$r_0$	$r_0$



Tree used for tree roles

# Tensor Product Decomposition Networks: Uncovering representations of structure learned by neural networks


R. Thomas McCoy,<sup>1</sup> Tal Linzen,<sup>1</sup> Ewan Dunbar,<sup>2</sup> and Paul Smolensky<sup>3,1</sup>

<sup>1</sup>Johns Hopkins University, <sup>2</sup>CNRS - Université Paris Diderot - Sorbonne Paris Cité, <sup>3</sup>Microsoft Research AI

# 5 Structure-Based Digit Sequence Tasks


- **RNNs trained to copy can be approximated almost perfectly:**

- Models being approximated: Sequential RNN and tree-based RNN trained to copy digit sequences.



- **Different tasks lead to different roles:**

- Reversal favors right-to-left where copying favors left-to-right
- With sorting, a non-structural task, bag-of-words roles work.




# 6 Sentence Encoder Experiments

- Test how often classifiers give the same output for a sentence encoding model and its tensor product approximation.

	Model Type	Training task
InferSent	BiLSTM	Natural Language Inference
Skip-thought	LSTM	Previous/next sentence prediction
SST	Tree	Sentiment prediction
SPINN	Tree	Natural Language Inference


- All 4 models are reasonably well approximated with non-structure-sensitive bag-of-words roles, suggesting they do not have robust representations of structure:



# 7 Related and ongoing work

- **Role learning (Soulos et al. 2019; arXiv:1910.09113):**

- Instead of using a role scheme generated by hand, add a module that automatically learns a role scheme.
- Analyzing a model trained on SCAN (Lake and Baroni 2018), a task of mapping a command to a sequence of actions:
  - jump twice → JUMP JUMP
  - walk after jump opposite right → RTURN RTURN JUMP WALK



- When applied to the sentence encoders, the role learner still does not outperform bag-of-words roles.

- **Using tensor product representations to solve tasks:**

- **Math problem solving:** Schlag et al. 2019; arXiv:1910.06611; Chen et al. 2019; arXiv:1910.02339
- **Question answering:** Palangi et al. 2017; arXiv:1705.08432
- **Image-caption generation:** Huang et al. 2017; arXiv:1709.09118

# 8 Conclusion

- 3 important puzzles about neural networks:
  1. How do neural networks represent structured information?
  2. How do they learn these representations?
  3. How do they use these representations to perform so well?
- Our work suggests an answer to puzzle 1: When trained on sufficiently structure-sensitive tasks, RNNs implicitly implement tensor product representations.
- Puzzles 2 and 3 remain for future work.

# Links and Acknowledgments

- Paper: <https://openreview.net/pdf?id=Bjx0sjC5FX>
- Demo: [http://rtmccoy.com/tpdn/tp\\_demo.html](http://rtmccoy.com/tpdn/tp_demo.html)
- This material is based upon work supported by the NSF GRFP, an NSF INSPiRE grant, an ERC grant (BOOTHON), and ANR grants IEC, PSL\*, GEOMPHON, USPC, and EFL. All opinions are our own.