

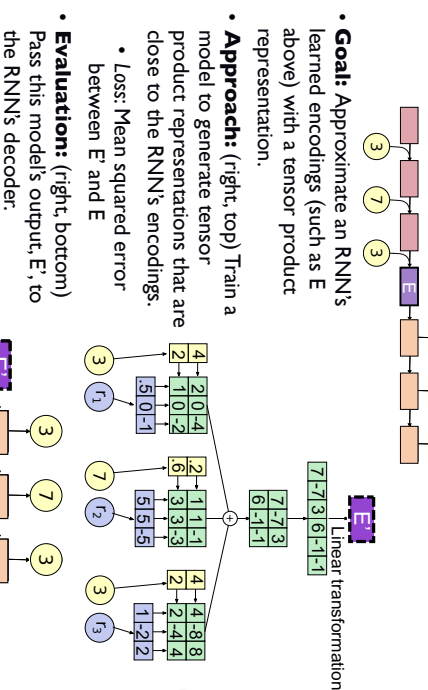
1 Overview

- **The puzzle:** How do recurrent neural networks (RNNs) use continuous vectors to represent discrete symbolic structures?
- **Findings:**
 - RNNs trained on structure-dependent tasks learn to implicitly implement *tensor product representations*.
 - Several popular tasks for training sentence encoders are not structure-sensitive enough to induce RNNs to capture sentence structure.

2 Tensor Product Representations


- A principled method for representing compositional symbolic structures in vector space (Smolensky 1990)
- Represent the input with pairs of **fillers** and **roles**:
Cats chase dogs = **cats:subject** + **chase:verb** + **dogs:object**
- Each filler f_i and role r_i has a vector embedding.
- The representation of the input is the sum of the outer products of each f_i and r_i : $\sum f_i \otimes r_i$

3 Tensor Product Decomposition




4 Role Schemes

	3	1	1	1	6
Left-to-right	0	1	2	3	
Right-to-left	3	2	1	0	
Bidirectional	(0,3)	(1,2)	(2,1)	(3,0)	
Wickelroles	#_1	3_1	1_6	1_#	
Tree	L	RLL	RRL	RR	
Bag-of-words	r0	r0	r0	r0	




5 Structure-Based Digit Sequence Tasks

- **RNNs trained to copy can be approximated almost perfectly:**
- Models being approximated: Sequential RNN and tree-based RNN trained to copy digit sequences.



- **Different tasks lead to different roles:**

- Reversal favors right-to-left where copying favors left-to-right
- With sorting, a non-structural task, bag-of-words roles work.




6 Sentence Encoder Experiments

- Test how often classifiers give the same output for a sentence encoding model and its tensor product approximation.

Model Type	Training task
Infer-Sent	BILSTM Natural Language Inference
Skip-thought	LSTM Previous/next sentence prediction
SST	Tree Sentiment prediction
SPINN	Tree Natural Language Inference

- All 4 models are reasonably well approximated with non-structure-sensitive bag-of-words roles, suggesting they do not have robust representations of structure:




Tensor Product Decomposition Networks:

R. Thomas McCoy,¹ Tal Linzen,¹ Ewan Dunbar,² and Paul Smolensky^{3,1}

¹Johns Hopkins University, ²CNRS - Université Paris Diderot, ³Microsoft Research AI

7 Related and ongoing work

- **Role learning** (Soulos et al. 2019; arXiv:1910.09113): Instead of using a role scheme generated by hand, add a module that automatically learns a role scheme.
- Analyzing a model trained on SCAN (Lake and Baroni 2018), a task of mapping a command to a sequence of actions:
 - jump twice → JUMP JUMP
 - walk after jump opposite right → RETURN RETURN JUMP WALK



- When applied to the sentence encoders, the role learner still does not outperform bag-of-words roles.

- **Using tensor product representations to solve tasks:**

- **Math problem solving:** Schlag et al. 2019; arXiv:1910.06611; Chen et al. 2019; arXiv:1910.02339
- **Question answering:** Palangi et al. 2017; arXiv:1705.08432
- **Image-caption generation:** Huang et al. 2017; arXiv:1709.09118

8 Conclusion

- 3 important puzzles about neural networks:
 1. How do neural networks represent structured information?
 2. How do they learn these representations?
 3. How do they use these representations to perform so well?
- Our work suggests an answer to puzzle 1: When trained on sufficiently structure-sensitive tasks, RNNs implicitly implement tensor product representations.
- Puzzles 2 and 3 remain for future work.

Links and Acknowledgments

- Paper: <https://openreview.net/pdf?id=Bjx0sIC5FX>
- Demo: http://rtmccoy.com/lpdr/lpdr_demo.html
- This material is based upon work supported by the NSF GRFP, an NSF INSPIRE grant, an ERC grant (BOOTHFON), and ANR grants IEC, FSL*, GEOPHON, USPC, and EFL. All opinions are our own.